

Temporal-Difference Learning

TD Prediction

- Both TD and Monte Carlo methods use experience to solve the prediction problem
 - TD methods learn from current predictions, and don't wait for actual returns
- TD(0) - One-step TD
 - Updates value function immediately after receiving state-value pair
- TD methods combine the sampling of Monte Carlo with the value bootstrapping of DP
- TD Error
 - Difference between the estimated value of S_t and the better estimate of $R_{t+1} + \gamma V(S_{t+1})$
- α is the learning rate

Monte Carlo Value Estimation

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

TD Value Estimation

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated
 Initialize $V(s)$ arbitrarily (e.g., $V(s) = 0$, for all $s \in \mathcal{S}^+$)
 Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 $A \leftarrow$ action given by π for S
 Take action A , observe R, S'
 $V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$
 $S \leftarrow S'$
 until S is terminal

Advantages of TD Prediction Methods


- TD methods update their estimates based in parts on other estimates (they bootstrap)
 - They learn a guess from a guess
- In practice, TD methods converge faster than MC methods
 - TD only needs to wait one time step to update its values - not an entire episode
 - TD methods are much more feasible than Monte Carlo methods for approximating solutions for problems with large state spaces

Optimality of TD(0)

- Batch Updating
 - Value function is only updated after processing a batch of training data (episode)
 - Monte Carlo methods find the estimates that minimize MSE on the training set
 - TD(0) finds the estimates that would be correct for the maximum-likelihood model (the model most likely to generate the data) of the Markov Process
- Certainty-Equivalence Estimate
 - Estimate of the value function is computed as if the model (observed from experience) was exactly correct - i.e. known with certainty
 - What TD(0) converges to, and why TD(0) converges faster than Monte Carlo

Sarsa: On-Policy TD Control

- Learns action-value function for the policy (epsilon-greedy) it is following
 - Learns from quintuples of: state-action reward \rightarrow state-action (Sarsa)
- TD(0) is used to estimate V , then epsilon-greedy policy improvement is used
- Sarsa methods learn *during the episode*, and therefore move out of poor policies quickly

Improvement Process of Sarsa**Sarsa Action-Value Update**

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
 Repeat (for each episode):
 Initialize S
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)
 Repeat (for each step of episode):
 Take action A , observe R, S'
 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$
 $S \leftarrow S'; A \leftarrow A'$
 until S is terminal

Q-learning: Off-Policy TD Control

- The learned action-value function directly approximates the optimal action-value function, independent of the policy being followed
 - Updates action-value function from the greedy policy
 - Chooses action from an epsilon-greedy policy
 - Converges slower than Sarsa, but can continue learning while changing policies

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
 Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)
 Take action A , observe R, S'
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 $S \leftarrow S'$
 until S is terminal

Expected Sarsa

- Like Q-Learning, but uses the expected value of state-action pairs, rather than the maximum of state-action pairs
- Moves deterministically in the same direction as Sarsa moves in expectation
- Expected Sarsa is more complex computationally than Sarsa, but it eliminates the variance due to the random selection of A_{t+1}
- Unlike Sarsa, Expected Sarsa can set $\mathbf{a} = \mathbf{1}$ without suffering any consequences, and can therefore obtain better short-term results
- If Expected Sarsa is used as an off-policy algorithm (exploratory behavior policy and greedy decision policy) than Expected Sarsa is exactly like Q-Learning
 - Expected Sarsa subsumes and generalizes Q-Learning while reliably improving over Sarsa - at the cost of a small additional computational cost

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Backup Diagrams for Q-Learning and Expected Sarsa

Q-Learning takes the maximum of action-value pairs (represented by the arc)

Expected Sarsa uses the expected value of action-value pairs



Maximization Bias and Double Learning

- Maximization Bias
 - Both Q-Learning and Sarsa use maximization to construct their target policies
 - This can lead to overestimating the Q value for actions with random rewards
- Double Learning
 - Two Q functions - Q_1 and Q_2 are independently learned
 - One function is used to determine the maximizing action, and the second to estimate its value - this removes maximization bias
 - Either Q_1 or Q_2 is updated randomly with the following equation

$$Q_1(S_t, A_t) \leftarrow Q_1(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q_2(S_{t+1}, \arg \max_a Q_1(S_{t+1}, a)) - Q_1(S_t, A_t) \right]$$

Double Q-learning

```

Initialize  $Q_1(s, a)$  and  $Q_2(s, a)$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily
Initialize  $Q_1(\text{terminal-state}, \cdot) = Q_2(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
    Choose  $A$  from  $S$  using policy derived from  $Q_1$  and  $Q_2$  (e.g.,  $\epsilon$ -greedy in  $Q_1 + Q_2$ )
    Take action  $A$ , observe  $R, S'$ 
    With 0.5 probability:
       $Q_1(S, A) \leftarrow Q_1(S, A) + \alpha (R + \gamma Q_2(S', \arg \max_a Q_1(S', a)) - Q_1(S, A))$ 
    else:
       $Q_2(S, A) \leftarrow Q_2(S, A) + \alpha (R + \gamma Q_1(S', \arg \max_a Q_2(S', a)) - Q_2(S, A))$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
  
```

Games and Afterstates

In games like tic-tac-toe, afterstate value functions would assess identical positions the same regardless of how the positions was reached, and the learning would transfer